

Daniel Filan

+1 (510) 646-5530 | Berkeley, CA | df@danielfilan.com | axrp.net | danielfilan.com

WORK EXPERIENCE

Member of Research Staff

Guidelight AI Standards

Jul 2026 —

Berkeley, CA

- At Guidelight, I write and assess safety standards for frontier AI companies
- One example of my work has been contributing to our [AI control assessment](#)

Member of Technical Staff

METR

Jan 2026 — May 2026

Berkeley, CA

- Project managed evaluation efforts leading into METR's [Frontier Risk Report](#)
- Was responsible for the compilation of the evaluation appendices of the same

Senior Research Manager

MATS Research, Inc.

May 2025 — Dec 2025

Berkeley, CA

- Continued with previous responsibilities, and additionally managed a research manager

Research Manager

MATS Research, Inc.

May 2024 — May 2025

Berkeley, CA

- Conducted a mix of personal and project management for researchers entering the fields of AI alignment, security and transparency (working with the researchers themselves and their mentors)
- Ran the process of selecting mentors for future cohorts

PODCAST

AI X-risk Research Podcast (AXRP) (axrp.net)

Dec 2020 — Present

- Host and publish a podcast featuring long-form interviews with researchers whose work focusses on analysing and/or reducing catastrophic and existential risk from artificial intelligence
- Funded by repeat grants from the [Long-Term Future Fund](#) and [Coefficient Giving](#)
- Scott Aaronson [wrote](#) of my interview with him: “The end result is ... well, probably closer to my current views on this subject than anything else I’ve said or written!”
- Stefan Schubert [tweeted](#) that “the episodes [of AXRP about AI policy] I’ve listened to have been excellent and epistemically fastidious”

EDUCATION

University of California, Berkeley

PhD (Computer Science)

Berkeley, CA

Aug 2016 — May 2024

- Thesis: “Structure and Representation in Neural Networks”, supervised by Stuart Russell

Australian National University

PhB (Hons)

Canberra, Australia

Feb 2012 — Dec 2015

- GPA 7.0/7.0, 1st class honours, University Medal
- Primarily studied mathematics and physics, honours in computer science
- Honours thesis: “Resource-bounded Complexity-based Priors for Agents”, supervised by Marcus Hutter

PUBLICATIONS

- **Constrained belief updates explain geometric structures in transformer representations.** Mateusz Piotrowski, Paul M. Riechers, *Daniel Filan*, Adam S. Shai. ICML, 2025.
- **Graphical clusterability and local specialization in deep neural networks.** Stephen Casper, Shlomi Hod, *Daniel Filan*, Cody Wild, Andrew Critch, Stuart Russell. PAIR²Struct Workshop, ICLR, 2022.
- **Exploring hierarchy-aware inverse reinforcement learning.** Chris Cundy, *Daniel Filan*. 1st Workshop on Goal Specifications for Reinforcement Learning, FAIM, 2018.
- **Self-modification of policy and utility function in rational agents.** Tom Everitt, *Daniel Filan*, Mayank Daswani, and Marcus Hutter. AGI, 2016.
- **Loss bounds and time complexity for speed priors.** *Daniel Filan*, Jan Leike, and Marcus Hutter. AISTATS, 2016.